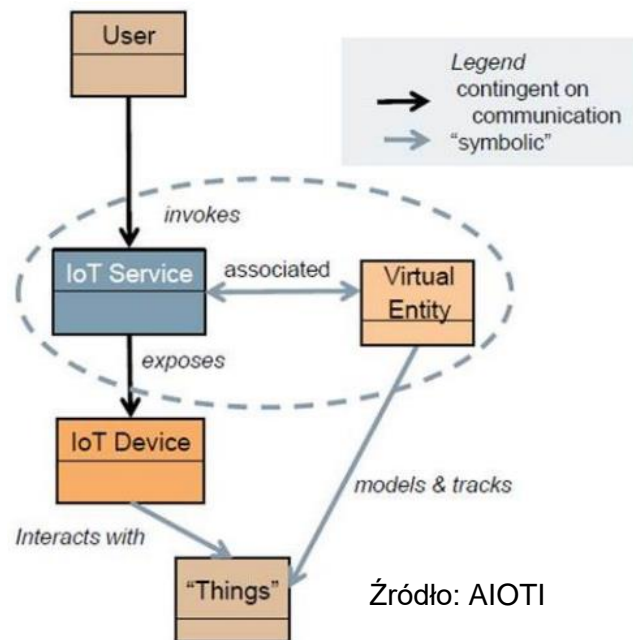


PSIR

Wprowadzenie do projektu 2024Z: gra Penney's ante

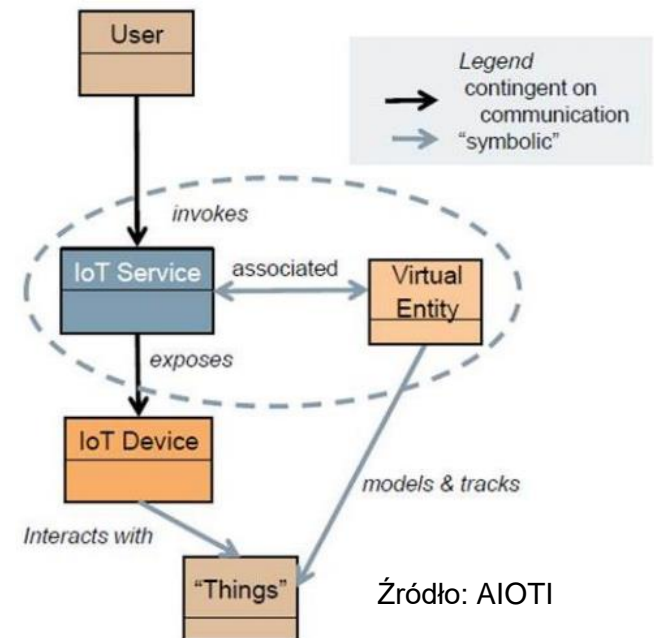
Jarosław Domaszewicz

Instytut Telekomunikacji Politechniki Warszawskiej



PSIR project concept and history

2

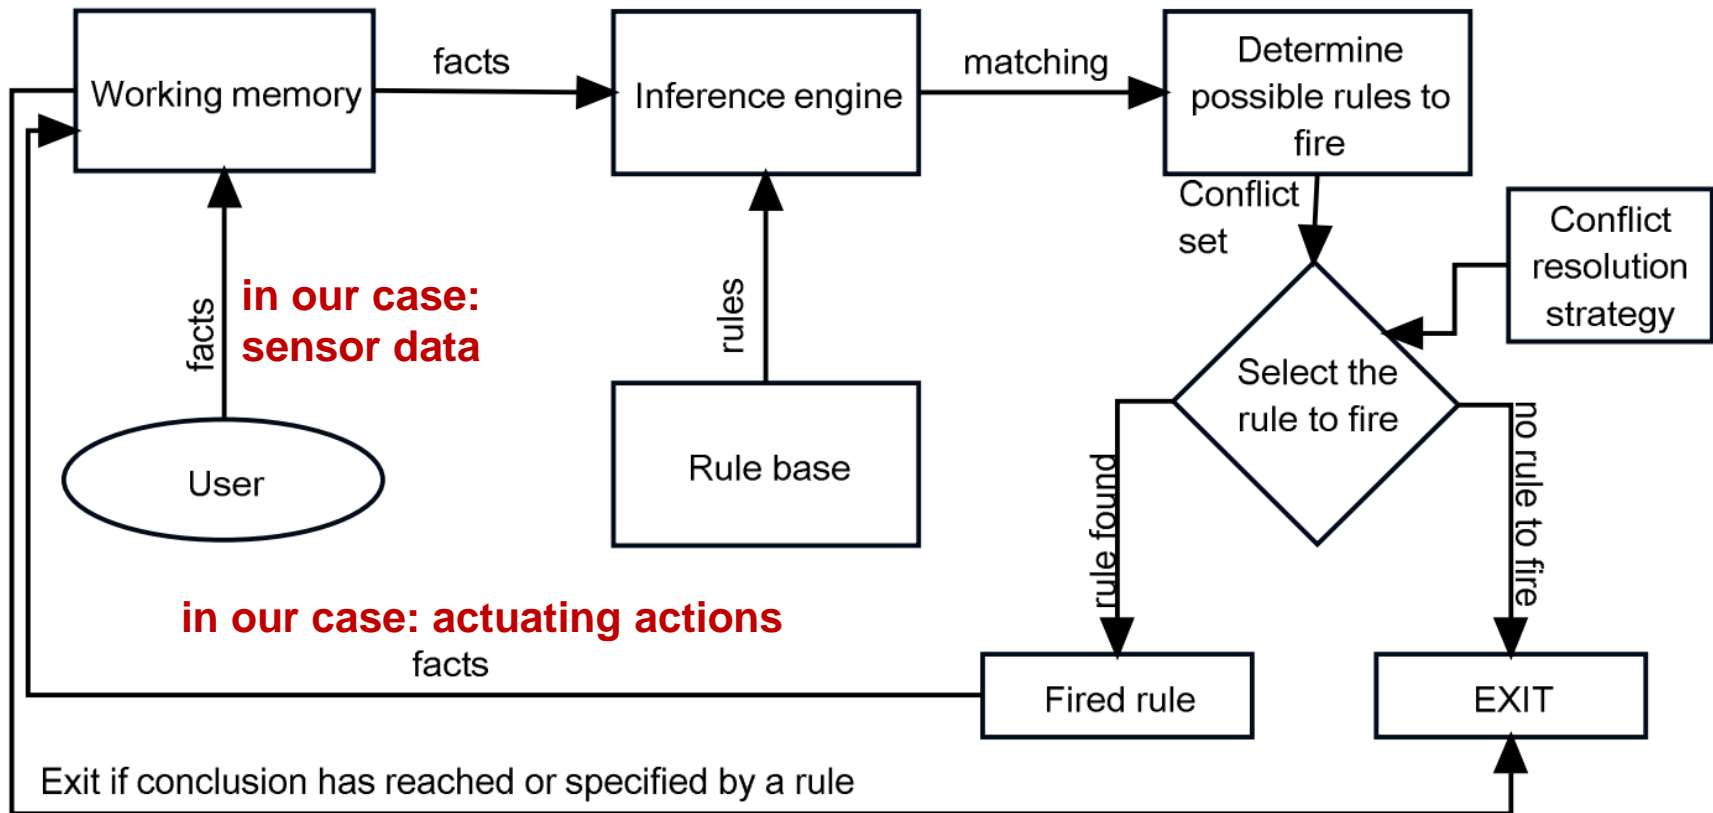


PSIR PROJECT CONCEPT

A programming challenge that helps develop solid programming skills while being intellectually stimulating.

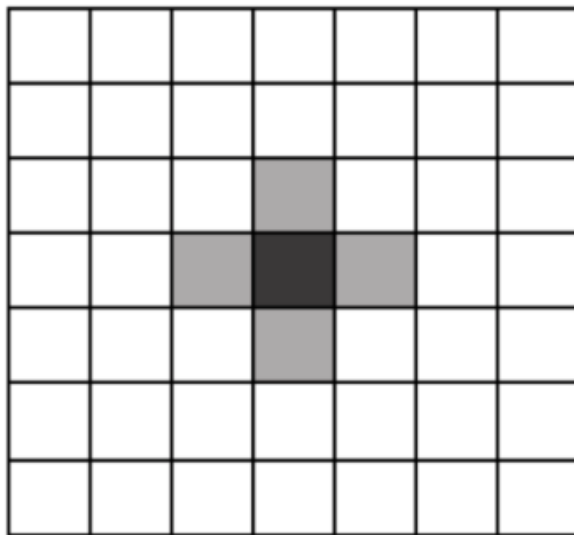
Enjoy stretching your mind!

21Z: A SMART ENVIRONMENT RULE-BASED SYSTEM

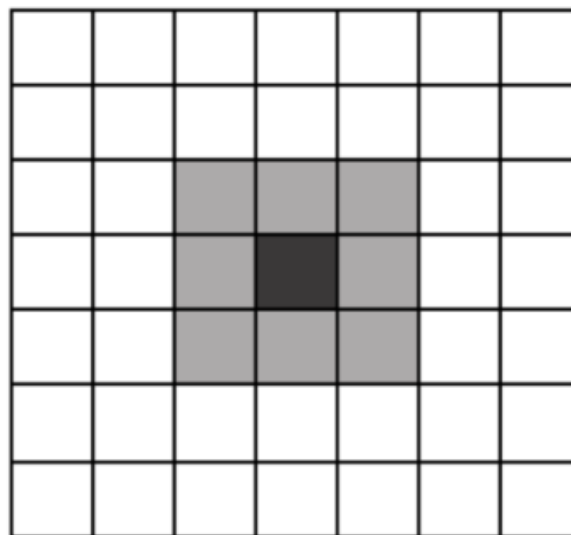


Source:
Intelligent Systems A Modern Approach,
Crina Grosan, Ajith Abraham,
Series: Intelligent Systems Reference Library, Springer,
DOI 10.1007/978-3-642-21004-4
Chapter 7, Rule-based Expert Systems, pp. 149-185

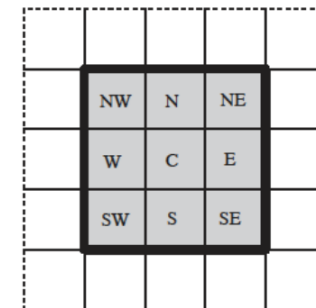
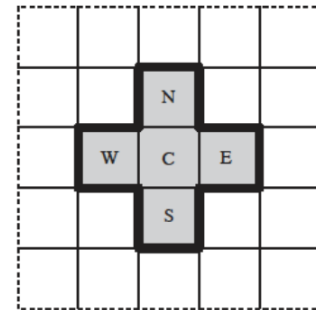
22L: CELLULAR AUTOMATA WITH SPACE PARTITIONING



Von Neumann



Moore

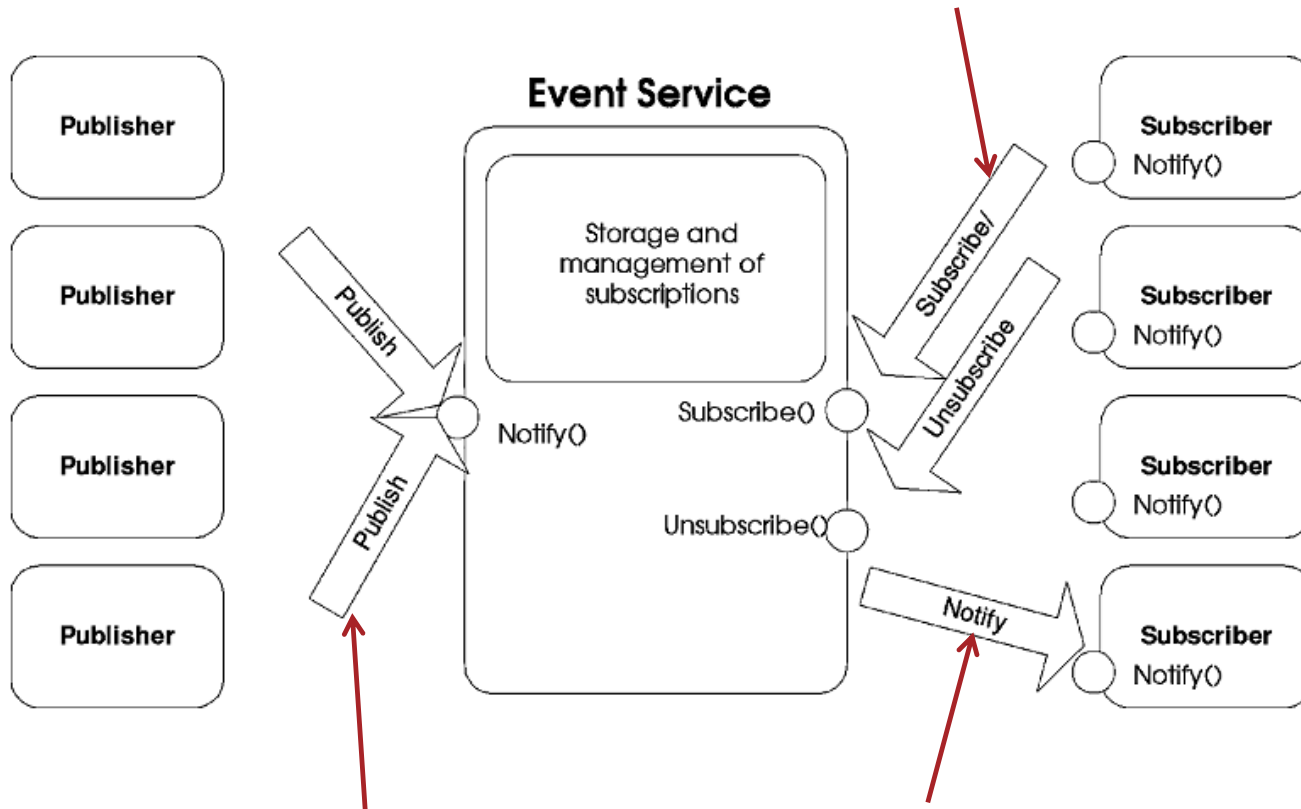


Source: Joseph Quartieri, Nikos E. Mastorakis, Gerardo Iannone, Claudio Guarnaccia
A Cellular Automata Model for Fire Spreading Prediction

Source: Washington Velasquez, Andres Munoz-Arcentalesy, Thomas Michael Bohnertz, Joaquin Salvachua
Wildfire Propagation Simulation Tool using Cellular Automata and GIS

22Z: CONTENT-BASED PUBLISH/SUBSCRIBE

1. subscribers express interest in selected events

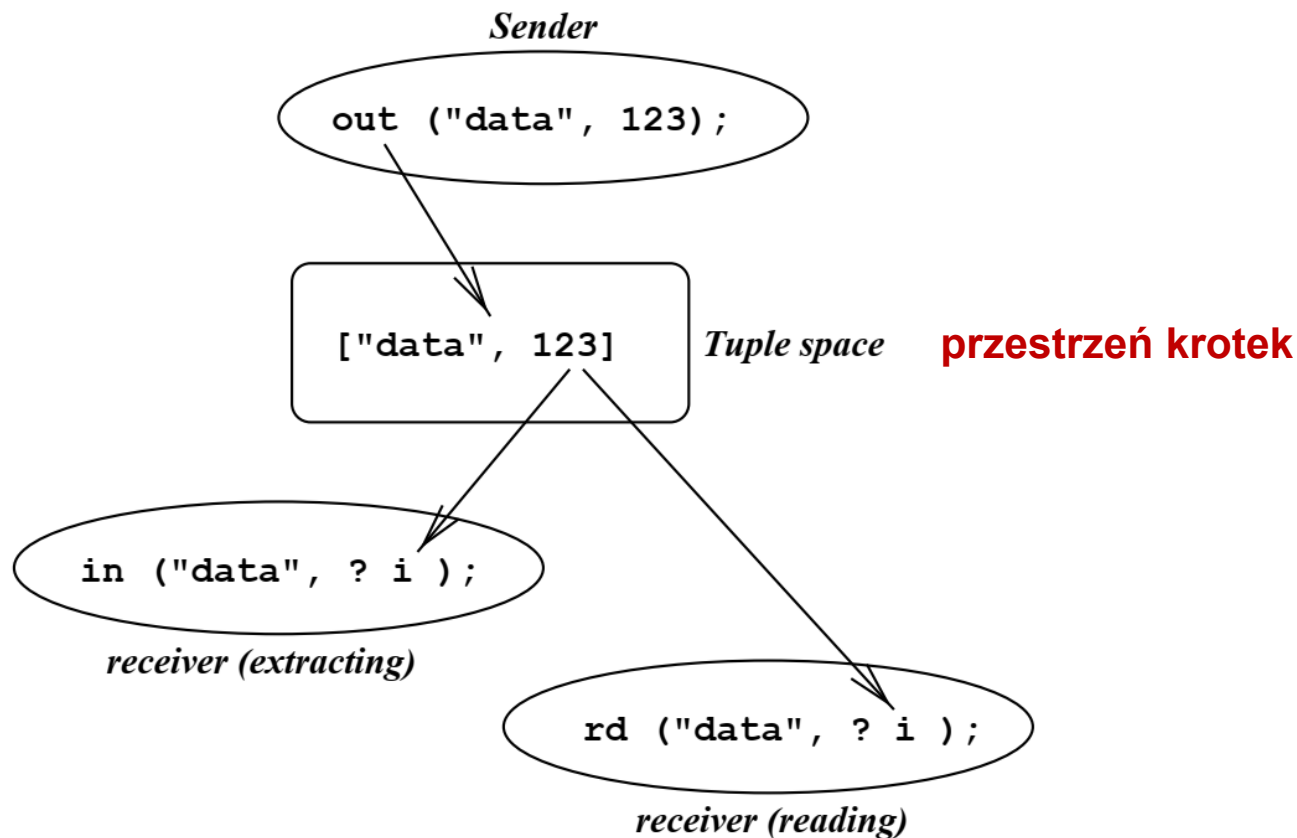


Source:
The Many Faces of Publish/Subscribe
P. Eugster et al.,
ACM Computing Surveys,
Vol. 35, No. 2, 2003

2. publishers produce assorted events 3. subscribers receive events they are interested in

- *publishers, subscribers, event service* (broker, server, middleware)
- note: push vs. pull, one-to-many, many-to-one

23Z: TUPLE SPACE MIDDLEWARE

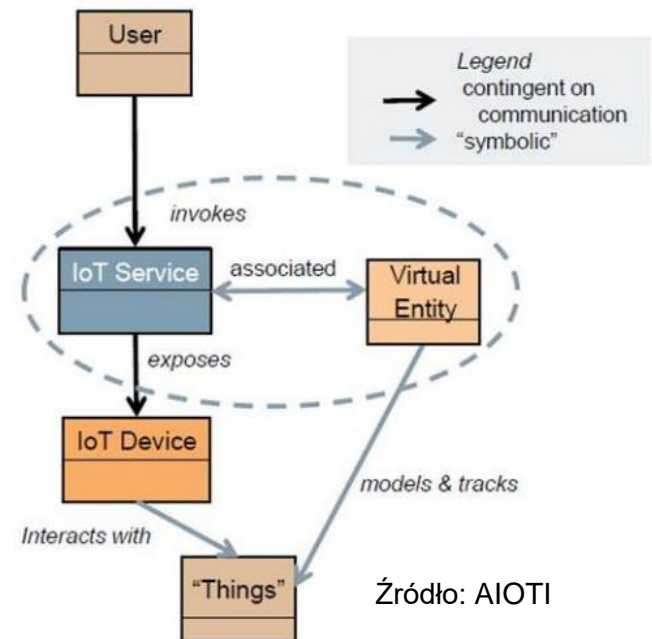


W. Hasselbring and M. Roantree. 1998.

A Generative Communication Service for Database Interoperability.

In Proceedings of the 3rd IFCIS International Conference on Cooperative Information Systems (COOPIS '98). IEEE Computer Society, USA, 64–73.

Penney's ante gane



NAVIA AUT CAPUT

Coin flipping, coin tossing, or heads or tails

is the practice of throwing a coin in the air
and checking which side is showing when it lands,
in order to randomly choose between two alternatives.

Coin flipping was known to the Romans as *navia aut caput* ("ship or head").



PENNEY'S ANTE GAME

Alice and Bob have a fair coin to flip.

They have decided to play a game known as **Penney's game**.

Alice selects a pattern of heads (H) and tails (T) of length n ,
after which Bob chooses his own pattern of heads and tails, also of length n .
They then begin tossing the coin.

Whoever's pattern appears first in the sequence of heads and tails is the winner.

Suppose $n = 3$, and Alice selects HHH, while Bob selects THH.

They toss the coin several times and get H H T T H T H H.

We see that Bob is the winner of this round.

SOME INTERESTING QUESTIONS

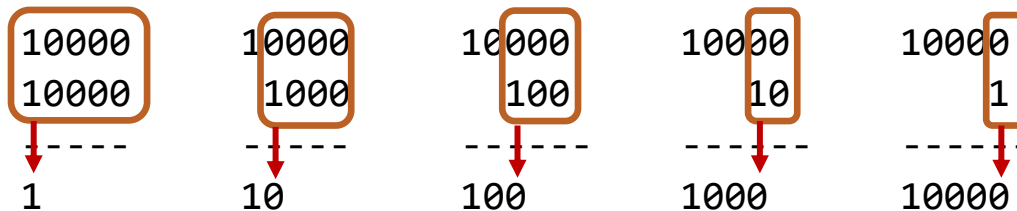
- Recall: Alice selects HHH, while Bob selects THH
- Are Bob and Alice equally likely to win?
- What is the probability that Alice wins?
- What is the probability that Bob wins?
- What is, on average, the number of coin tosses needed to complete a game?
- What about different combinations of patterns and different lengths of patterns (n)?
- What if there are more than two players?

A SIMPLE EXAMPLE

- Alice picks HH, and Bob picks TH.
- The probability of each pattern is $\frac{1}{4}$, so the game appears fair.
- Is it?
- Consider the results of the first two tosses.
 - HH... -> Alice wins
 - TH... -> Bob wins
 - HT... -> can Alice win?
 - TT... -> can Alice win?
- Probability of Alice winning is $\frac{1}{4}$.
- Probability of Bob winning is $\frac{3}{4}$.

MAGIC OF CONWAY LEADING NUMBERS (1/3)

- Let w_1 and w_2 be two patterns.
- Define a Conway leading number as in the following example.
- $w_1 = 10000$, $w_2 = 11111$
- C^{w_1, w_1}



- Keep shifting the second argument with respect to the first argument and checking if the prefix of the second argument is the same as the suffix of the first argument.
- Let's calculate C^{w_1, w_2} , C^{w_2, w_1} , and C^{w_2, w_2}

MAGIC OF CONWAY LEADING NUMBERS (2/3)

- Let P_1 and P_2 be the probabilities that w_1 and w_2 occurs first, respectively. Then

$$\frac{P_1}{P_2} = \frac{C^{2,2} - C^{2,1}}{C^{1,1} - C^{1,2}}$$

- In our example ($w_1 = 10000$, $w_2 = 11111$):

$$\frac{P_1}{P_2} = \frac{31-1}{16-0} = \frac{30}{16} = \frac{15}{8}$$

- You can derive general formulas for probabilities in two-player Penney's ante.
 - use the above formula for the ratio of probabilities and $P_1 + P_2 = 1$
 - the formulas express the probabilities in terms of the Conway leading numbers

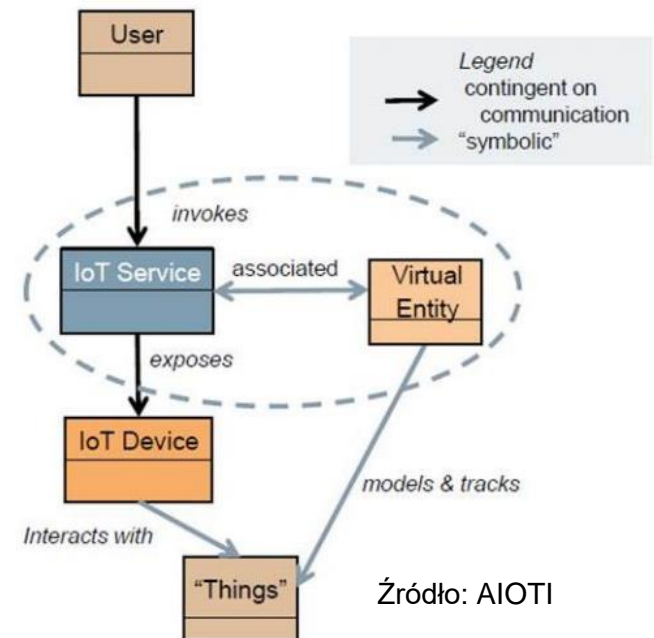
MAGIC OF CONWAY LEADING NUMBERS (3/3)

- One can also derive a formula for the expected value of the number of tosses before any of the two patterns occurs.

$$E = 2 \times \frac{C^{1,1}C^{2,2} - C^{1,2}C^{2,1}}{C^{1,1} - C^{1,2} + C^{2,2} - C^{2,1}}$$

Our objective: distributed Penney's ante

16

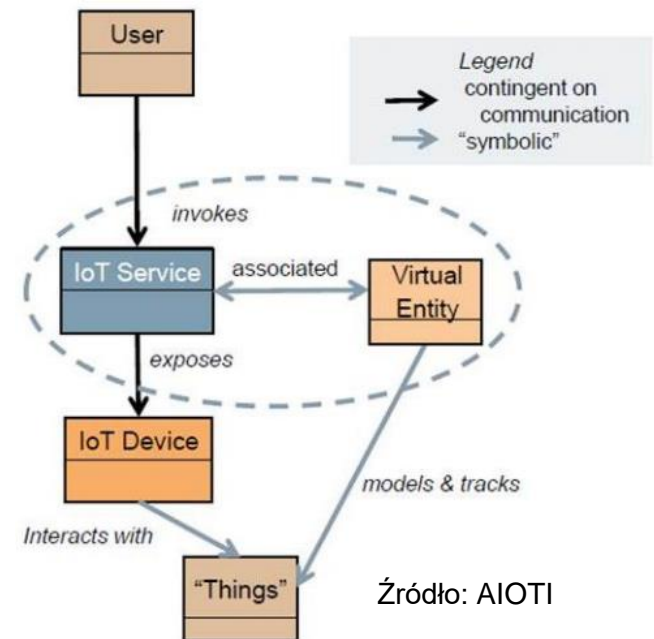


OUR OBJECTIVE

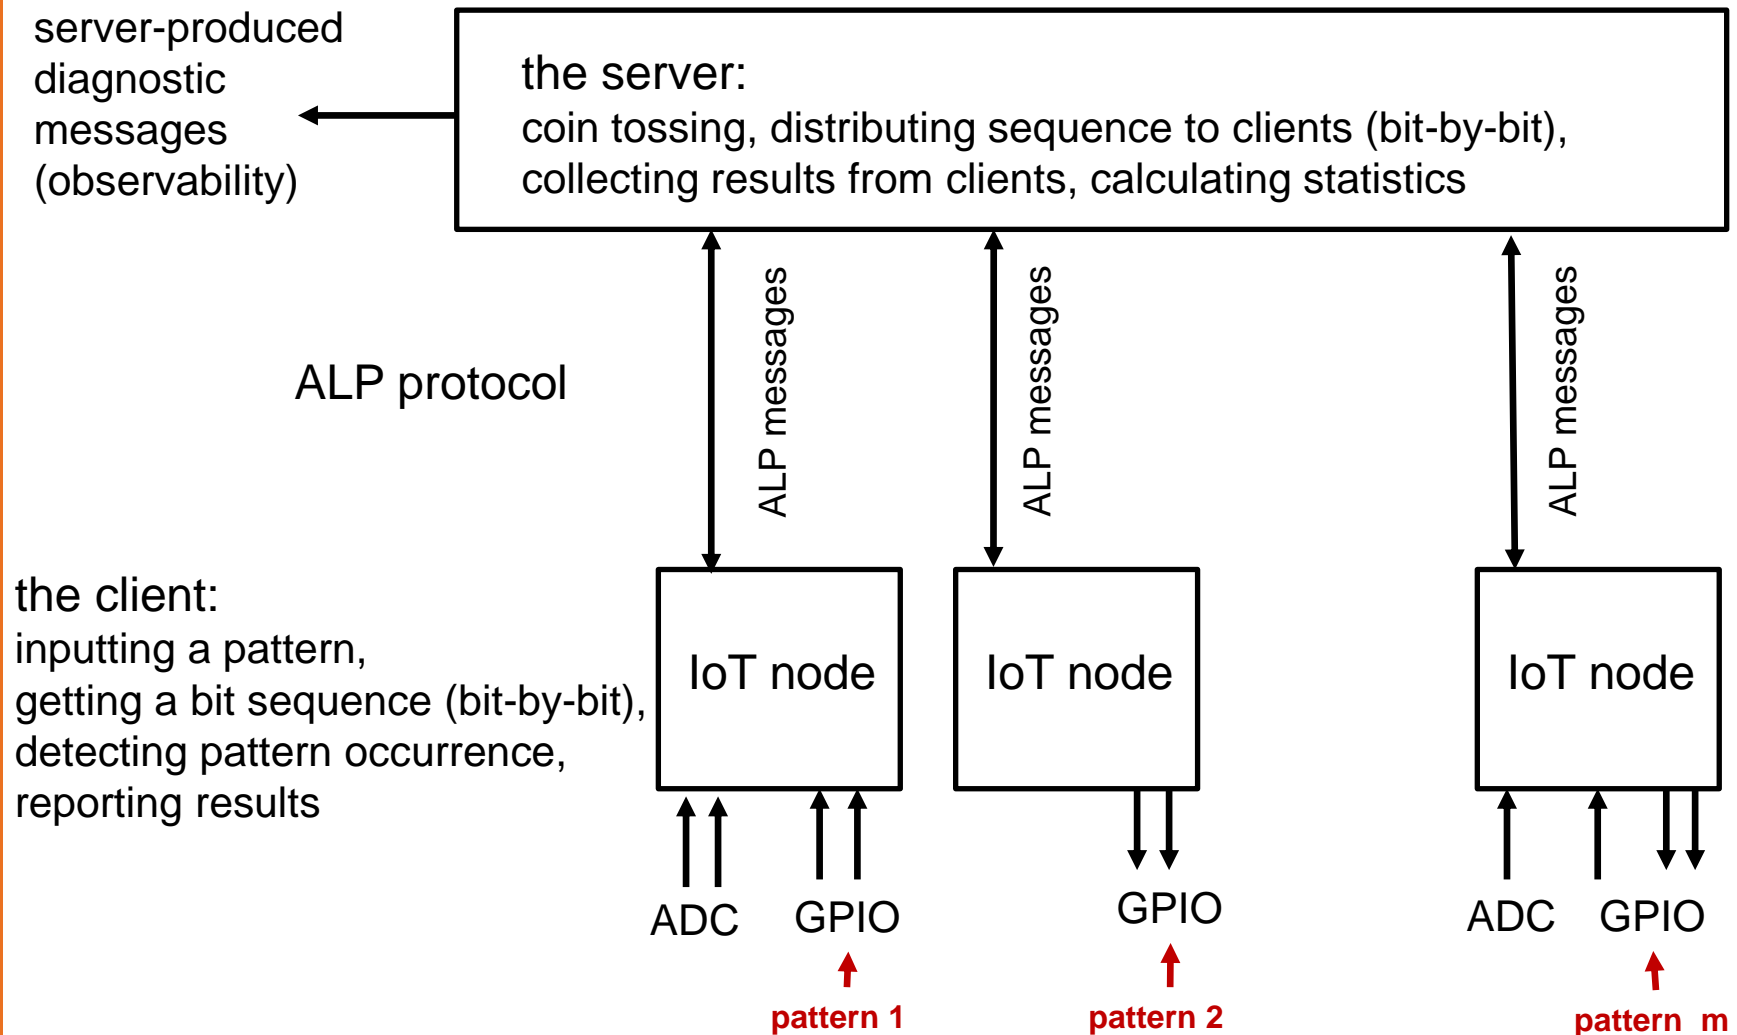
- We'll make a distributed application that simulates the Penney's ante game.
 - not much computation, but a lot **communication and coordination**
- For a given set of patterns, we'll play the game multiple times.
- Major results from each game played:
 - which pattern is the winner
 - how many tosses were needed for the winning pattern to occur
- Major results from multiple games (played with same patterns):
 - for each pattern, (an estimate of) the **probability that the pattern wins**
 - (an estimate of) the **average number of tosses** needed for a winning pattern to occur

System high-level architecture

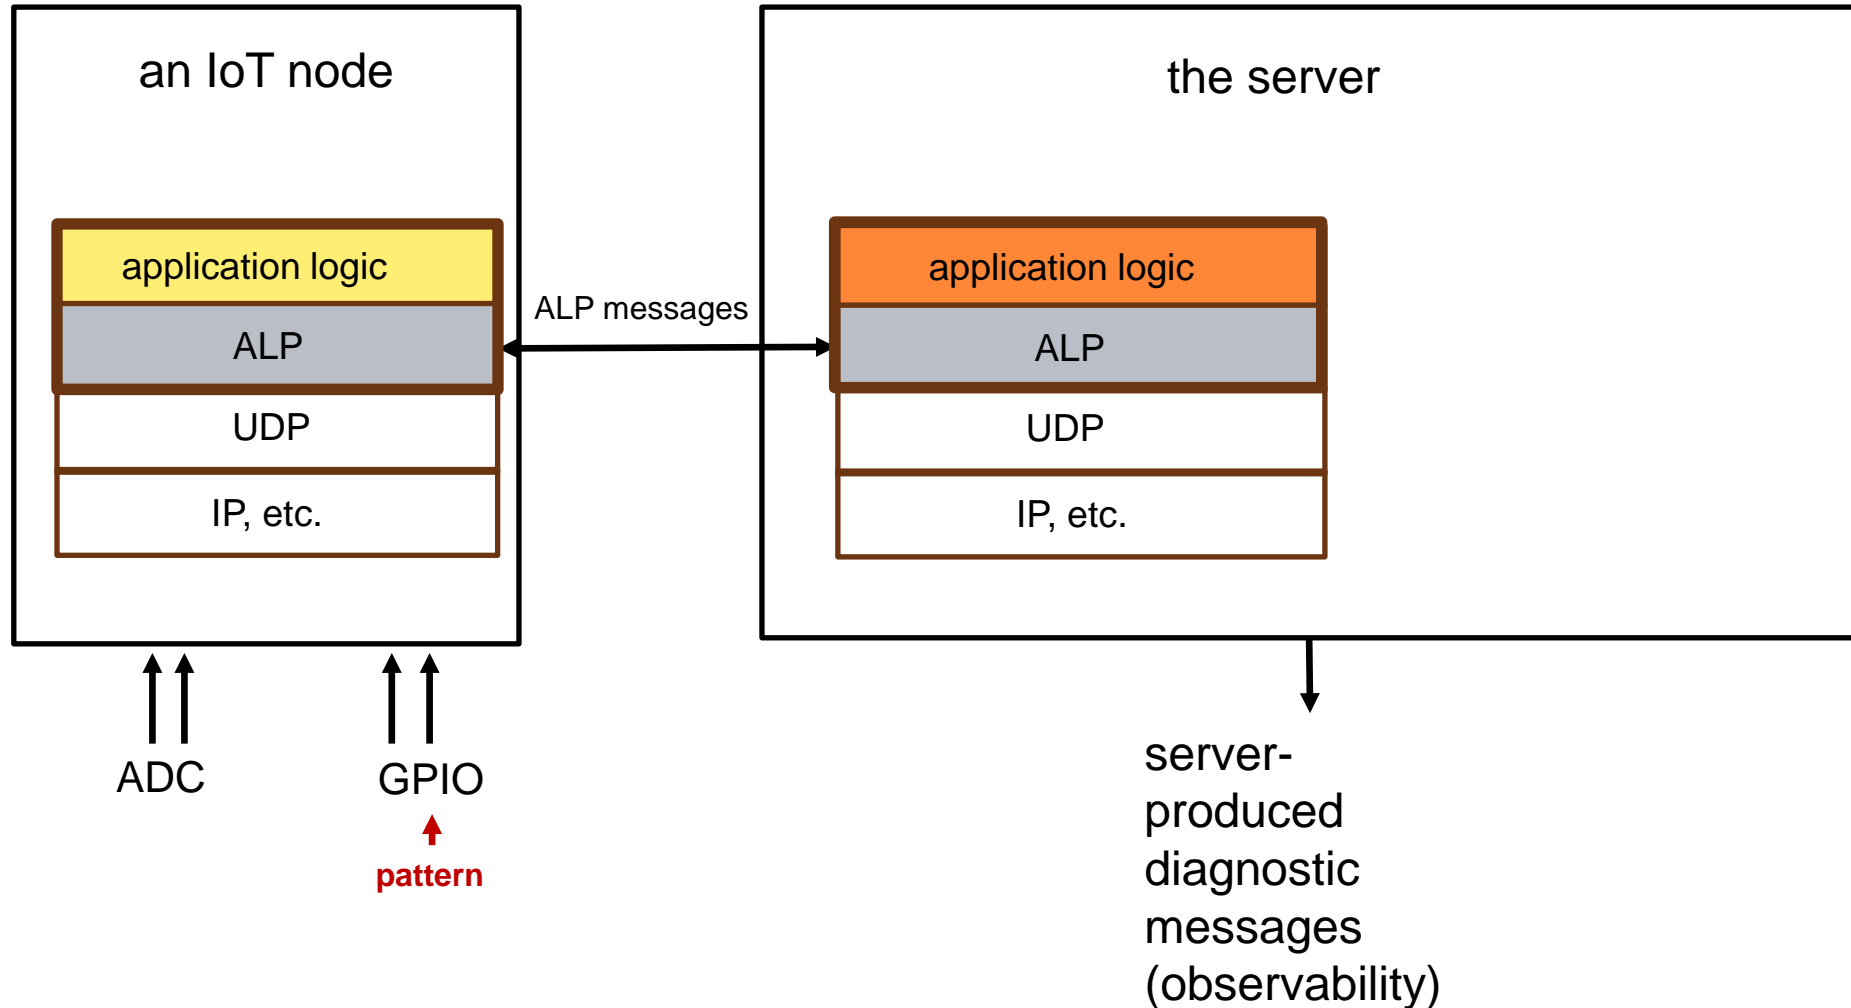
18



SYSTEM ARCHITECTURE

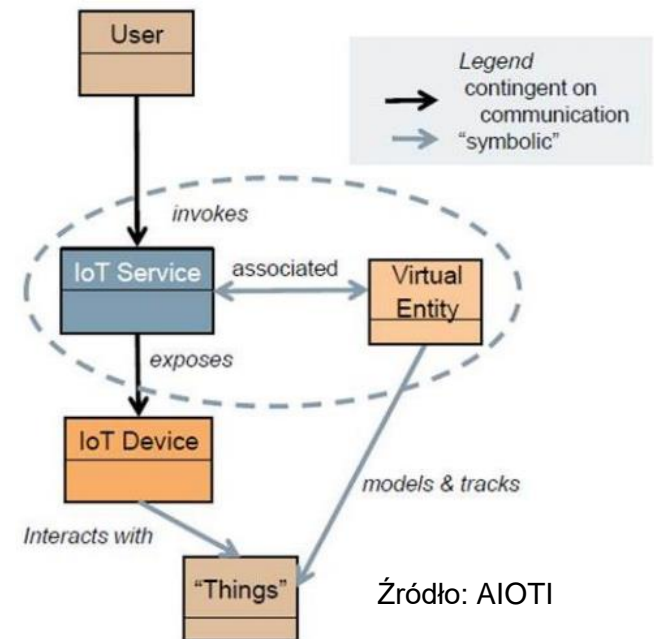


SYSTEM ARCHITECTURE



ALP protocol

21



ALP

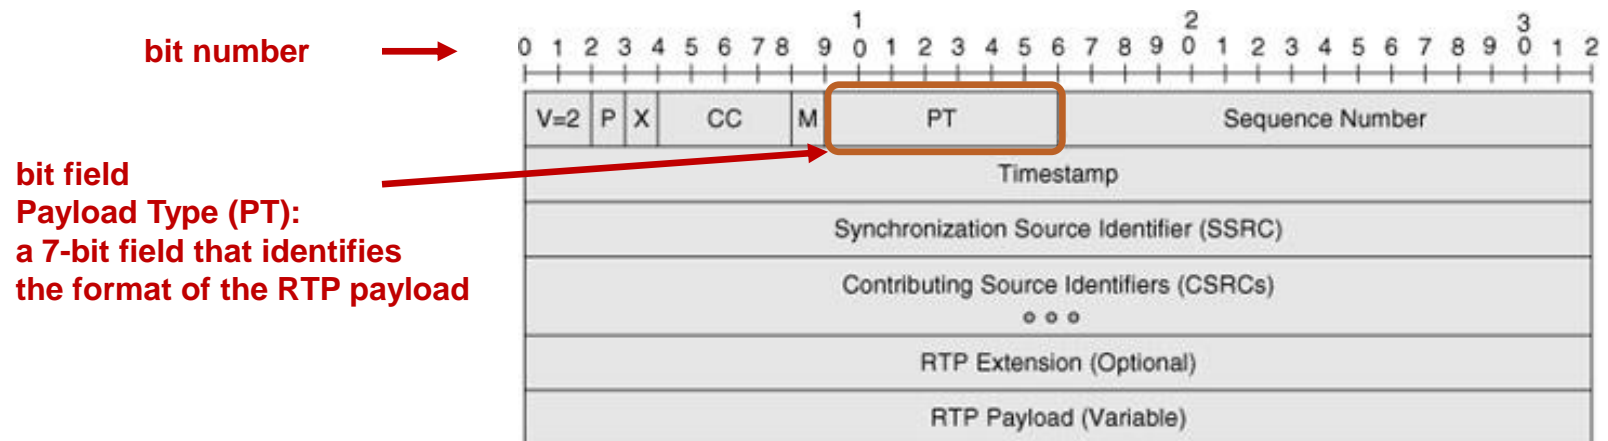
- An application layer protocol.
- Used for communication between a client and a server.
- It's entirely your creation.
- It should be language and platform independent.
 - e.g., clients based on different platforms, written in different languages may participate in the distributed application
 - btw, this is nothing new: making it possible for participating parties to be heterogeneous is the benefit of having a well specified protocol (and one of the main reasons to specify it)

ALP IS A BINARY PROTOCOL (1/2)

- „zaprzyjaźnij się z bajtem” (*befriend the byte*)
- ALP should be a *binary* protocol
- why? to save memory and the amount of transmitted data
- you need to define message formats at the bit level
 - identify bit fields within words
 - for each bit field, specify encodings (meanings of different bit patterns)

ALP IS A BINARY PROTOCOL (2/2)

- example: RTP (Real-time Transport Protocol)
 - a binary application layer protocol to transfer, e.g., voice in „phone calls“ (VoIP)



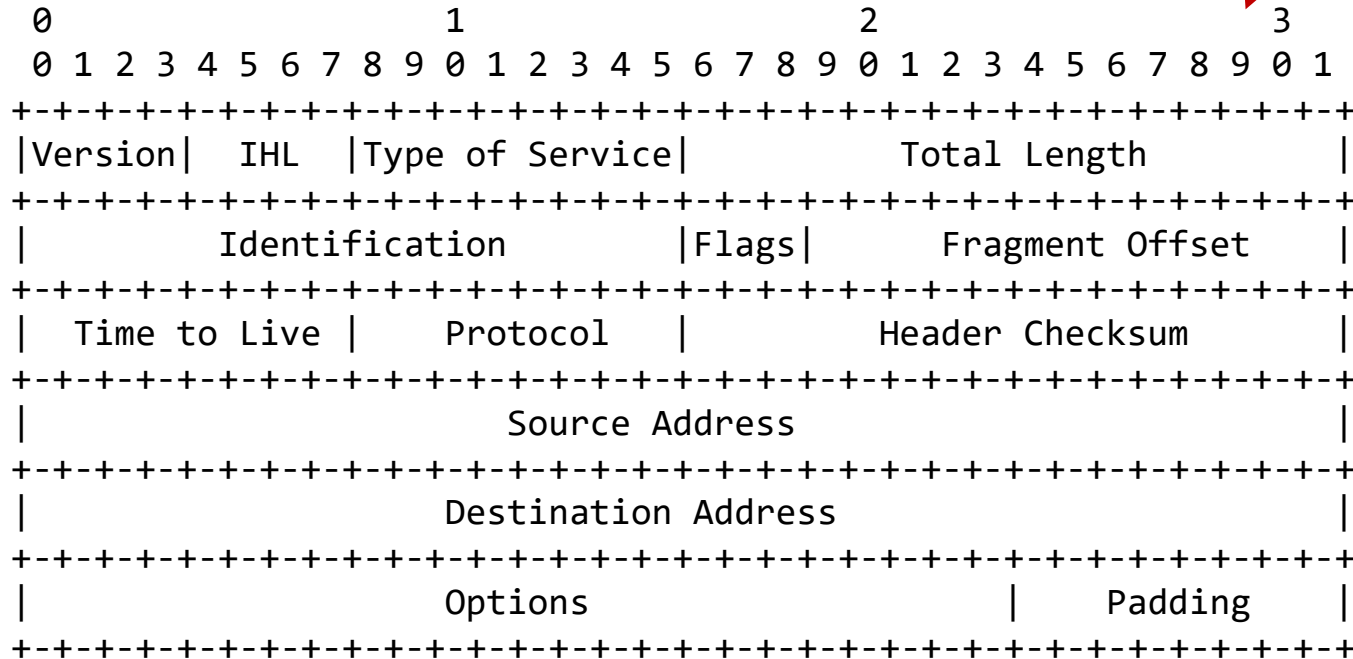
- you may also look at RFC 791 (Fig. 4) or RFC 793 (Fig. 3)
- hint: when working with a binary protocol, make sure you understand things like:
 - network byte order
 - endianness

ALP IS A BINARY PROTOCOL (3/3)

document the format
of ALP messages like this



- another example: IPv4 (RFC 791)



Example Internet Datagram Header

Figure 4.

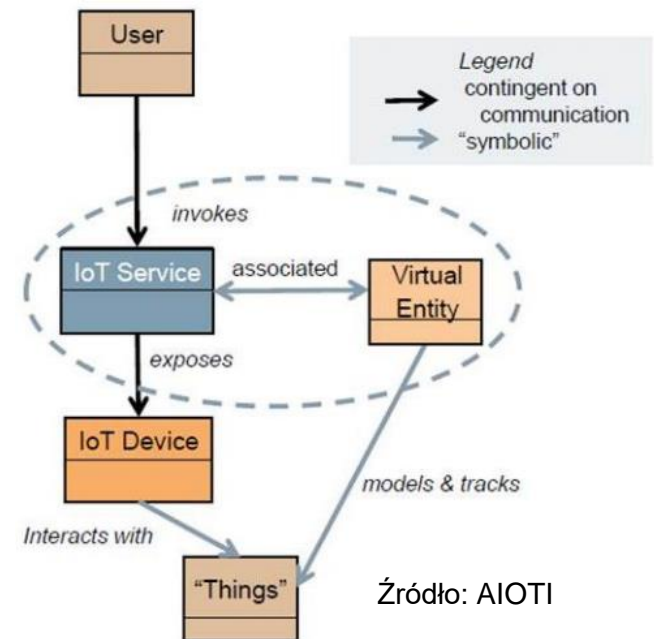
- in the case of ALP, you also need to specify the format of the payload
- you may also look at RFC 793 (Fig. 3)

ALP SHOULD RUN ON TOP OF UDP

- UDP often used in IoT due to its simplicity
- unreliable
 - add a simple reliability scheme
 - transmit, wait for ACK, if no ACK, retransmit

Server

27



SERVER FUNCTIONALITY 1

- Collects registrations from participating nodes.
 - this way finds out how many nodes will play

SERVER FUNCTIONALITY 2

- In a loop, manages a given number of Penney's ante games:

```
1. for(i=0; i<num_games; i++)  
2.     make_one_game();
```

SERVER FUNCTIONALITY 3

- One game looks as follows:

```
1. make_one_game() {
2.
3.     start a game by sending a start message to all the nodes;
4.
5.     do {
6.         toss a coin (use a random number generator);
7.         // the coin should be fair (both probabilities equal to ½)
8.         distribute H or T (one bit) to all the nodes;
9.     }
10.    while (no node reports that its pattern has occurred);
11.
12.    // due to the lack of synchronization, reports from multiple nodes may arrive
13.    // ... but only one node is the winner ...
14.    // ... namely the one that reports the least number of tosses
15.    // some reports (e.g., the winning one) may be delayed with respect to others
16.    // make sure you receive all reports
17.
18.    make a record for the game just ended;
19.    // record which pattern won
20.    // record how many tosses were needed for the winning pattern to occur
21. }
```

SERVER FUNCTIONALITY 4

- Once all the games are played, calculate estimates:
 - for each pattern, the probability that the pattern wins
 - the average number of tosses needed for a winning pattern to occur
- (It suffices to take averages.)
- Display info about the games:
 - the number of players
 - the length of the patterns (n)
 - how many times the game was played
 - the probabilities of winning (estimate)
 - the average number of tosses (estimate)

SERVER FUNCTIONALITY 5

- Important: do not re-initialize your random number generator after each game.
 - otherwise, you may end up playing with the same tossing history many times

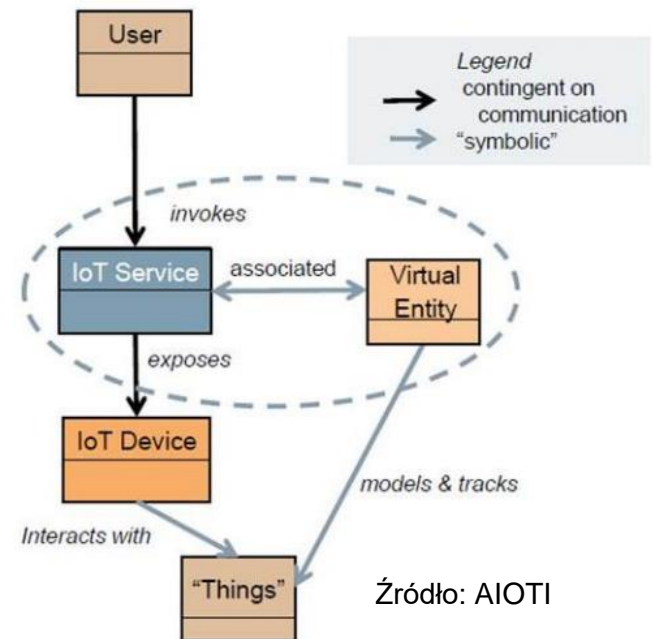
SERVER FUNCTIONALITY 6

- While doing all of the above, keep updating metrics and produce diagnostic messages.
 - examples:
 - the number of all messages sent and received so far
 - average message length
 - the number of completed games
 - ...

SERVER PLATFORM

- Use Linux system API and C language facilities only
 - no libraries, no platforms, no middleware, ...

Nodes



NODE FUNCTIONALITY

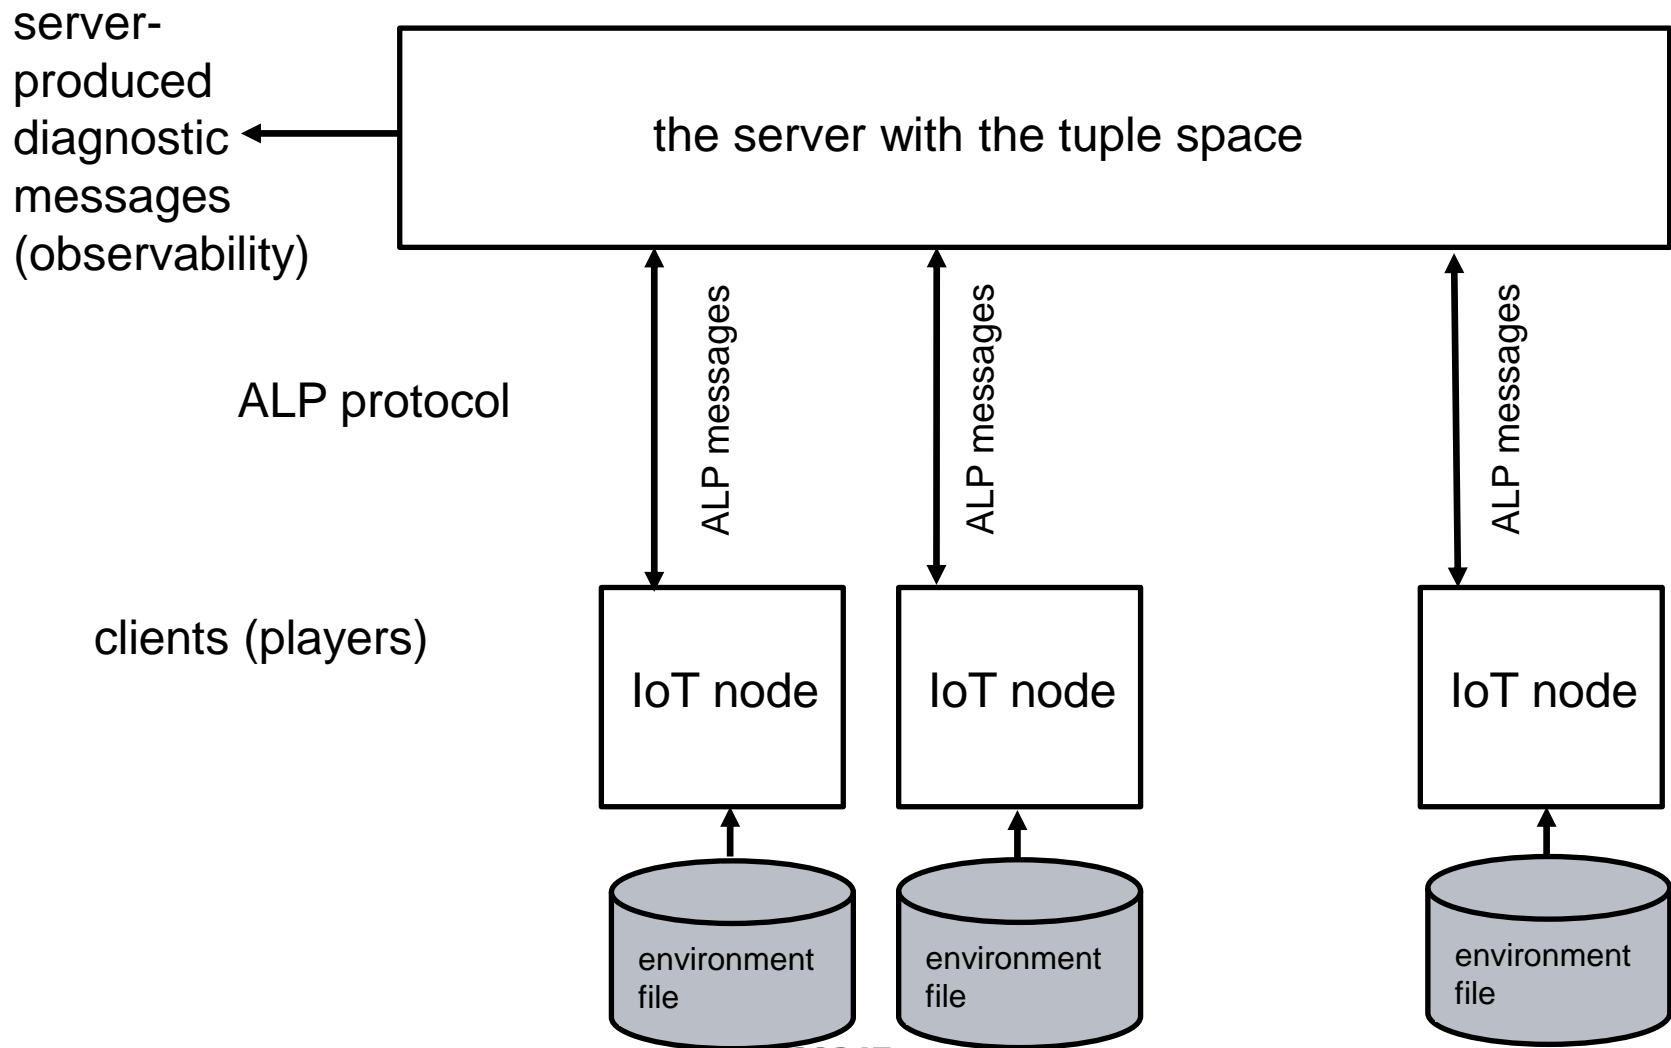
- Reads the node's pattern, using GPIO.
- Registers with the server.
- In a loop, plays Penney's ante games, as directed by the server:

```
1. do {
2.     wait_for_message;
3.     if(start) {
4.         clear_tossing_history; //important! each new game should start from scratch
5.         continue;
6.     }
7.
8.     // make sure the message carries the result of a toss (H or T)
9.     update tossing history;
10.    if(pattern has occurred)
11.        send a report to the server
12.    }
13. } while (1);
```

NODE PLATFORM

- Arduino emulator
- Use "standard" Arduino API and PSIR extensions only.
- The emulator supports UDP only.

ENVIRONMENT FILES



ENVIRONMENT FILE

```
+ qTemperature,quantity,Z0      # input 0=-10 1023=+20
+ qHumidity,quantity,Z1
+ sOpening,status,D1           # input 0=OPEN, 1=CLOSED
+ aSwitch,action,D2            # output 0=ON, 1=OFF
```

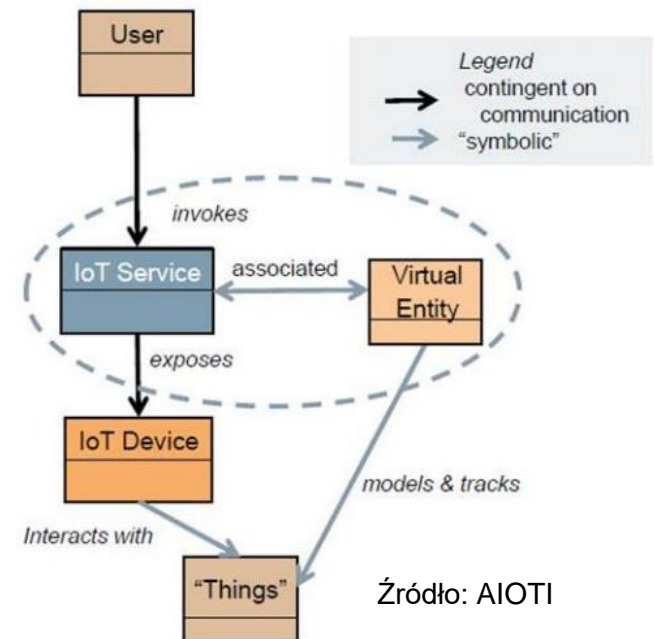
```
: 1000,qAirTemperature, 20
: 3000,qAirTemperature, 21
: 5000,sOpening, 0
: 7000,qHumidity, 512
```

HOW TO INPUT A PATTERN

- The length of the pattern (n) can be hard-coded.
- Use two GPIO pins.
- To get one bit of a pattern:
 - detect an edge on one GPIO pin
 - right after detecting the edge, read the value on the other GPIO pin
- Repeat the above until you get n bits.

Your results

41



YOUR RESULTS

- source code
- a report
- a demonstration

YOUR RESULTS ARE YOURS!!!

- Freely talk with other students about concepts.
- You may reuse some items developed by another team, as long as you give credit to the team that created them.
 - this is a general rule that helps you avoid plagiarism
 - however, we reserve the right to make a judgment as to how important the reused items are, and to deduct points accordingly
 - reusing somebody else's work may turn out very costly in terms of points deducted
- We can easily see reusability!
- If we see reused items without clearly given credit, all teams with those items (including the authors) will have the same number of points deducted.
 - (we cannot investigate who has taken what from whom)
- Your best strategy, if you are cooperative: offer others advice but do not share items you are going to submit as results!

THE CONTENTS OF YOUR REPORT

1. a specification of your ALP

- produce message formats and binary encodings for the protocol messages
 - when specifying message formats, follow a good example, e.g., RTP or IP
- produce sequence diagrams (see the UML language to learn about sequence diagrams)
- describe how you deal with the fact that UDP is unreliable

2. a description of your server implementation

- components (overview of the server architecture)
- major data structures

3. a description of your node implementation

- include a description of your environment files

4. estimates obtained for some sets of patterns

1. probability of winning for each pattern
2. average number of tosses

YOUR DEMO

- prepare several nodes (players)
 - each node requires its own environment file
- during the demo
 - start your server
 - start all the players
 - ...
 - wait for estimates

GRADING

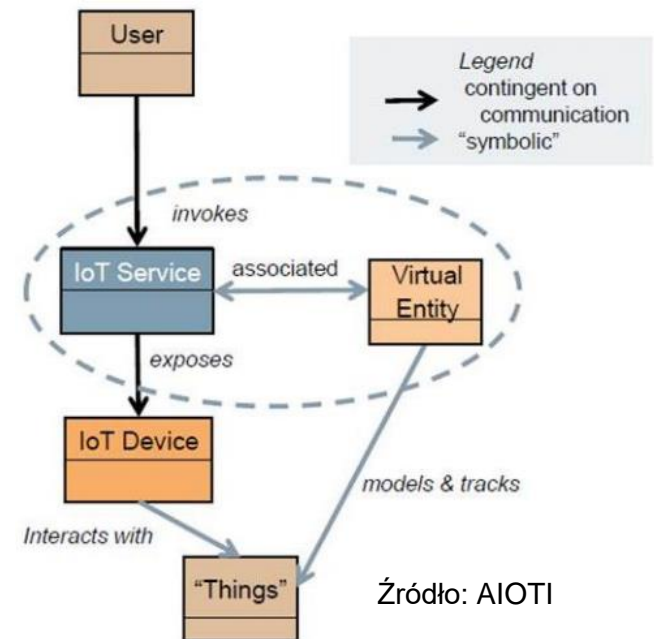
- source code quality 10%
- report 45%
 - 1-20%, 2-10%, 3-10%, 4-5%
 - (see the contents of the report)
- demo 45%
 - what counts most is whether your software works ...
 - ... but whether the demo is presented smoothly counts as well

DEADLINE

- All teams must upload their results by January 20, 2025 12:00.
 - that's Monday
- Uploading to the PSIR-supplied git account (also used for labs).

References

48



REFERENCES

- Isha Agarwal, Matvey Borodin, Aidan Duncan, Kaylee Ji, Tanya Khovanova, Shane Lee, Boyan Litchev, Anshul Rastogi, Garima Rastogi, Andrew Zhao *From Unequal Chance to a Coin Game Dance: Variants of Penney's Game* <https://doi.org/10.48550/arXiv.2006.13002>
- Stanley Collings *Coin sequence probabilities and paradoxes* Bulletin of the Institute of Mathematics and its Applications (1982) 18, 227-232

Dziękujemy za uwagę!

50

